

An Exclusive

1. The recipient should use XOR to decrypt the ciphertext. When the ciphertext is XOR'd with the same one-time pad used to encrypt the message, the recipient will get the original message.

$$\begin{array}{r} 11100010 \quad 00011100 \\ \wedge 10101010 \quad 01010101 \\ \hline 01001000 \quad 01001001 \end{array}$$

2. The original text can be obtained by XOR'ing the ciphertext and the one-time pad as follows:

$$\begin{array}{r} 01100101 \quad 10100010 \quad 01101111 \\ \wedge 00111100 \quad 11100111 \quad 00111100 \\ \hline 01011001 \quad 01000101 \quad 01010011 \\ 89 \quad 69 \quad 83 \\ \text{ASCII} \quad Y \quad E \quad S \end{array}$$

Therefore, the three ASCII characters are Y, E, and S.

3. We can encrypt the longer message by using a key, which is generated by repeatedly concatenating the original key of 16 bits. For example, if the message is 64 bits long, then our new key will be generated by concatenating the original key of 16 bits 4 times. If the message is not a multiple of 16, for example if it is 72 bits long, then our new key will be generated by concatenating the original key 4 times and its left most 8 bits to produce a new key of 72 bits.
4. A longer key is more secure, because a ciphertext generated using a shorter XOR key can be decrypted using frequency analysis of characters. Since the frequency distribution of letters cannot be effectively masked using a shorter XOR key, using the typical distribution of English letters, one can break such ciphertexts. In addition, when a short XOR key is repeatedly used to encrypt a longer message, if one part of the message can be guessed, then the remaining part of the message can be easily broken.
5. The given function switches the values at the addresses pointed to by **a** and **b**. In other words, if the initial values of ***a** and ***b** are 1 and 2 respectively, the new values of ***a** and ***b** after the function is called will be 2 and 1 respectively. A better name for this function would be "swap".
6. If two out of the three drives fail, then we cannot recover those two drives. In this case, we only have one good drive.

7. We only need one additional hard drive. We need to XOR the data from the first three drives and store the results onto the fourth drive. For example, if the data in the first three drives are namely A, B, and C, then we will store $D = A \oplus B \oplus C$ onto the fourth drive (backup drive). Supposing one of the three drives fails, we can XOR the remaining two drives with our backup drive to recover the lost drive. This is because, since $D = A \oplus B \oplus C$, we know that $A = D \oplus B \oplus C$, $B = D \oplus A \oplus C$, and $C = D \oplus A \oplus B$.

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?
Wikipedia (XOR cipher article)
2. 90 minutes

Assembly Line

1. The function of each line in the program is listed below. The numbers correspond to the line numbers in the program:
 1. This statement sets the value of register `r1` to string `"x: "`.
 2. Displays the value of the contents of register `r1`. This will display `"x: "` on the screen.
 3. This statement prompts the user for an input value and when the user inputs a value, stores that value in register `r2`.
 4. This statement sets the value of register `r1` to string `"y: "`.
 5. Displays the value currently stored in register `r1`. In other words, it will display `"y: "` on the screen.
 6. This statement prompts the user for an input value. When the user inputs a value, it then stores that value in register `r3`.
 7. This statement evaluates whether the contents of the registers `r2` and `r3` are equal. If they are equal, the program jumps to line 11.
 8. Note that we reach this line only if the `JUMPEQ` evaluation in Line 7 fails. In this case, the content of the register `r1` is set to the string, `"x is not equal to y"`.
 9. This statement displays the current content stored in register `r1`. In other words, it displays `"x is not equal to y"`.
 10. The program exits execution.
 11. The program will only reach this line if the `JUMPEQ` statement on Line 7 evaluates true. In other words, the program will reach this point only if `x` is equal to `y`. This statement sets the value of register, `r1`, to the string `"x is equal to y"`.
 12. This statement prints the contents of the register `r1`. Therefore, the string `"x is equal to y"` is displayed.
 13. The program exits execution.

2. The following program, when the two numbers are not equal, instead of just printing out “x is not equal to y”, it prints either “x is less than y” or “x is greater than y”, depending on which number is greater. It still prints “x is equal to y” if the two numbers are equal.
 1. SET r1 “x: “
 2. PRINT r1
 3. INPUT r2
 4. SET r1 “y: “
 5. PRINT r1
 6. INPUT r3
 7. JUMPEQ r2 r3 15
 8. JUMPLT r2 r3 12
 9. SET r1 “x is greater than y”
 10. PRINT r1
 11. EXIT
 12. SET r1 “x is less than y”
 13. PRINT r1
 14. EXIT
 15. SET r1 “x is equal to y”
 16. PRINT r1
 17. EXIT
3. The following program “coughs” some (n) number of times input by the user.
 1. SET r1 “n: “
 2. PRINT r1
 3. INPUT r2
 4. SET r3 0
 5. SET r1 1
 6. SET r4 “cough”
 7. PRINT r4
 8. ADD r3 r3 r1
 9. JUMPLT r3 r2 6
 10. EXIT
4. The name of the instruction for calling a function is **callq**. This is because in the original C program, we are calling the function **get_int()** twice, and in the assembly code, we see that the instruction **callq** is followed by **get_int()** twice. Similarly, in the original C program, there are two calls to the **printf()** function, and we note that the assembly code has two occurrences of **callq** followed by **printf()**. This indicates that **callq** is the x86-64 assembly instruction for calling a function.
5. The name of the x86-64 instruction that determines what to print is **jge**, which means “jump if greater than or equal to”. The preceding instruction to **jge** in the assembly code is **cmpl**, which compares the contents (values of **x** and **y**) of the two registers **rbp** and **eax** and sets one of the bits in condition register field depending on the values of **x** and **y**. Based on the conditional bit, **jge** performs a comparison jump if the destination operand is greater than or equal to the source operand. From the assembly code, we note

that if the call to **jge** is successful, it jumps to line number 33 and prints the string **L.str.3** (which is, “**x is not less than y**”). (If the conditional jump fails, we print the string **L.str.2** (which is, “**x is less than y**”).

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

X64 Cheat Sheet from the CS department of Brown University

2. About how long, in minutes, did you spend on this problem's questions?

60 minutes

Bit by Bitcoin

1. A blockchain is a connected list of transactional records (or blocks), in which each block is linked to the previous block using a hash value of the previous block. In addition, each block also contains its own transactional data and a timestamp when it was added to the chain. The most important property of a blockchain is that it is resistant to tampering of data contained in any of the blocks. The concept of blockchain is analogous to a public transaction ledger that is used to record transaction between two parties. However, every transaction recorded in a blockchain is verified before it becomes permanent part of the chain. In other words, a blockchain is a public, digital online ledger. Once the transaction is recorded in a block, its contents cannot be altered without modifying all the blocks that were added subsequently.

The blockchain technology forms the foundation of the concept of a digital currency called the bitcoin. The bitcoin is a decentralized digital network that allows two parties to conduct monetary transactions directly without requiring an intermediary agency or a middleman (such as a traditional bank). The idea behind bitcoin is, in a manner similar to the physical currency, a person cannot spend more money than what she has. Since bitcoin network uses the blockchain technology, it is a shared public ledger, where all of its confirmed transactions are securely stored.

2. A cryptographic hash function is a mapping function, which maps an input data into a deterministic hash value. The most important property of a cryptographic hash function is that a hash value generated by the algorithm is practically impossible to invert. If the original message is known, then it is easy to verify whether the hash value corresponds to the message, but one cannot identify the message from the hash value. A good cryptographic function generates vastly different hash values even if the original message is only slightly changed.
3. Assuming the string s consists of ASCII characters, the given hash function will not change the ordinal values of the first 50 characters as the modulo operation yields the same ordinal value as the original character. The ordinal values of the remaining characters in the ASCII tables can be easily guessed by just adding a multiple of 50 to the hash value. Therefore, the given implementation of hash is unsuitable as a cryptographic hash function.

4. The following is an implementation of the block type as specified:

```
typedef struct
{
    int transaction_id;
    int sender_id;
    int recipient_id;
    float transaction_amount;
    digest sender_signature;
    digest proof_of_work;
}
block;
```

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

The video on bitcoin from www.3blue1brown.com

Wikipedia article on blockchain

Wikipedia article on bitcoin

2. About how long, in minutes, did you spend on this problem's questions?

80 minutes

Count on It

1. We can represent $2^7 = 128$ total characters in ASCII, if each character is represented using 7 bits.
2. The first byte read will tell us how many bytes the character corresponding to that byte occupies. The first byte for a character represented in UTF-8 encoding will have the following format depending on the number of bytes used for that character:

1 st Byte	Number of bytes used for the character
0xxxxxxx	1
110xxxxx	2
1110xxxx	3
11110xxx	4

Therefore, depending on the format of the first byte read, we need to further read either 0, 1, 2 or 3 more bytes for a single character. This process repeats until we finish reading the entire file.

3. If the first two bits of a byte are 10, then it is a continuation of an existing character. Therefore, if we XOR the byte with 10000000 (256) and obtain a value less than or equal to 64 (i.e: **00xxxxxx**), then the byte is a continuation of an existing character. In other words, a continuation byte is of the form **10xxxxxx**.
4. The following program counts the number of characters in a file that is encoded as UTF-8.

From my answer to Part (2) above, we note that the first byte of a character that is encoded in UTF-8 will be either less than 128 or greater than $128+64 = 192$. Therefore, the only check we need to do is to verify that the value of current byte is either less than 128 or greater than or equal to 192.

Please see the modified program on the next page.


```

#include <stdbool.h>
#include <stdio.h>

typedef unsigned char BYTE;

int main(int argc, char *argv[])
{
    if (argc != 2)
    {
        printf("Usage: ./count INPUT\n");
        return 1;
    }

    FILE *file = fopen(argv[1], "r");
    if (!file)
    {
        printf("Could not open file.\n");
        return 1;
    }

    int count = 0;
    while (true)
    {
        BYTE b;
        fread(&b, 1, 1, file);
        if (feof(file))
        {
            break;
        }

        // This may be a UTF-8 encoded file. Therefore, count
        // only the first byte of a character
        //
        if (b < 128 || b >= 192)
        {
            count++;
        }
    }

    printf("Number of characters: %i\n", count);
}

```

Be sure to submit Question 4 via submit50 as well!
I have submitted my code for Question 4 via submit50.

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?
 - a. Wikipedia article on UTF-8
 - b. www.fileformat.info
2. About how long, in minutes, did you spend on this problem's questions?

60 minutes

File Browser

1.

```
int fgetpos(FILE *stream)
{
    // The difference between current_byte and first byte, cast to long,
    // will be equal to how many bytes have been so far read.
    //
    return (long)(stream->current_byte - stream->first_byte);
}
```

2.

```
bool fsetpos(FILE *stream, int offset)
{
    // Make sure that we don't set the current byte beyond the end of the file
    if ((stream->current_byte - stream->first_byte + offset > stream->size) ||
        // Also make sure that a negative offset does not set the beginning of the
        // file to before stream->first_byte
        (stream->current_byte + offset < stream->first_byte))
    {
        printf("Error: Invalid offset\n");
        return false;
    }
    else
    {
        stream->current_byte = stream->first_byte + offset;
        return true;
    }
}
```

3.

```
typedef unsigned char BYTE;

bool fwrite(BYTE *buffer, int size, FILE *stream)
{
    // If the file is currently set to read_only = true, then it was
    // not opened with "w" option. Raise an error and return false.
    if (stream->read_only)
    {
        printf("Error: Cannot write to file\n");
        Return false;
    }
    // Copy the contents of length = size pointed to by buffer onto stream
    //
    memcpy(stream->current_byte, buffer, size);

    // Adjust the current_byte value
    stream->current_byte = stream->current_byte + size;

    // Update the file size
    stream->size = stream->size + size;

    Return true;
}
```

4.

```
bool feof(FILE *stream)
{
    // If the difference between the current_byte and first_byte exceeds
    // the size of the file, then the caller has read past the end of the
    // file.
    if ((stream->current_byte - stream->first_byte) > stream->size)
    {
        return true;
    }

    return false;
}
```

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

1. www.sourceware.org
2. File I/O in C programming with examples, www.beginnersbook.com

2. 90 minutes

Hey, Earl

1. **`^yes+$`** matches with any string starting with **"ye"** followed by one or more **"s"**, and ending with an **"s"**.
2. **`^yes\s*$`** matches with any string starting with **"yes"** followed by zero or more white spaces and ending in a white space.
3. We can modify **`re.compile`** as follows to stop the matching at the first white space. This will match up to the first name, so the output in the example will be always **"Hey, Earl"**:

```
p = re.compile("my name is ([^\s]+)", re.IGNORECASE)
```

4. The given DFA implements the regular expression, **`^$ | (EMMA)+`**
5. The given DFA will accept the following three strings:
 - a. **`""`** (empty string)
 - b. **`EMMA`**
 - c. **`EMMAEMMA`** (or any number of repetitions of **`EMMA`**)
6. An example of a string that the given DFA would not accept is **`CAT`**.

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

www.docs.python.org

2. About how long, in minutes, did you spend on this problem's questions?

120 minutes

Key Questions

1. Assuming the records for both **Max** and **Ileana** already exist in table users and their usernames are "**max**" and "**ileana**" respectively, we need to execute the following SQL statement to insert a row into the followers table:

```
INSERT INTO followers(follower_id, followee_id)
  SELECT VIEW1.id1, VIEW2.id2
    FROM (SELECT id AS id1
          FROM users
          WHERE username = "max") VIEW1,
         (SELECT id AS id2
          FROM users
          WHERE username = "ileana") VIEW2;
```

Please note that in the above query, we could have also used the **name** column (instead of the **username** column) in the WHERE clause of the subqueries (example: **name** = "**Max**" instead of **username** = "**max**").

2. The following query selects the usernames of every user that follows Reese and whom Reese also follows back:

```
SELECT DISTINCT username FROM users, followers,
  (SELECT followers.follower_id AS stalker_id
   FROM users, followers
   WHERE users.username = "reese" AND
         users.id = followers.followee_id) VIEW1
WHERE users.id = followee_id AND
      Followers.followee_id = VIEW1.stalker_id;
```

3. The following query returns the usernames of users who are "two degrees of separation" away from Reese.

```
SELECT DISTINCT users.username FROM users, followers,
  (SELECT DISTINCT followers.followee_id AS friends_id
   FROM users, followers
   WHERE users.username = "reese" AND
         users.id = followers.follower_id) VIEW1
WHERE VIEW1.friends_id = followers.follower_id AND
      followers.followee_id = users.id AND
      users.username != "reese";
```

4. The following SQL statement creates a **friendships** table that allows for the representation of friendships and friend requests. Please note that the column “**friend**” is of type INTEGER and can take the values of either 0 or 1, meaning “not a friend” or “friend”. (Sqlite does not have a boolean type, which would have been a more appropriate choice for this column.)

```
CREATE TABLE friendships
(
    user_id INTEGER,
    friend_requester_id INTEGER,
    friend INTEGER,
    FOREIGN KEY (user_id) REFERENCES users(id),
    FOREIGN KEY (friend_requester_id) REFERENCES users(id)
);
```

5. The following SQL statement inserts a row into the **friendships** table when Max sends a friend request to Ileana. Please note that this row is inserted with an initial value of 0 (zero) for the friend column. The initial value of the friend column is 0 indicates that the user specified by **user_id** has received a friend request from the user specified by **friend_requester_id**. (but the request has yet been accepted).

```
INSERT INTO friendships(user_id,
                        friend_requester_id,
                        friend)
SELECT VIEW1.id1,
       VIEW2.id2,
       0
FROM
    (SELECT id AS id1 FROM users
     WHERE username = "max") VIEW1,
    (SELECT id AS id2 FROM users
     WHERE username = "ileana") VIEW2;
```

6. The following SQL statement updates the value of the **friend** column in the friendships table for Ileana when she accepts Max’s friend request.

```
UPDATE friendships SET friend = 1
WHERE user_id IN
    (SELECT id from users WHERE username = "ileana")
AND
    friend_requester_id IN
    (SELECT id from users WHERE username = "max");
```

7. The following three SQL statements deletes all traces of Reese, his friendships, and his followees from the database. Please note that is it important to delete from the **users** table at the end, as the first two queries are dependent on being able to find **Reese** in the **users** table. Please also note that the following queries could also use the **name** column in the **users** table, instead of the **username** column. In this case, the queries will need to use "**Reese**", instead of "**reese**".

```
DELETE FROM followers
  WHERE follower_id IN
    (SELECT id FROM users
      WHERE username = "reese") OR
    followee_id IN
    (SELECT id FROM users
      WHERE username = "reese");
```

```
DELETE FROM friendships
  WHERE user_id IN
    (SELECT id FROM users
      WHERE username = "reese") OR
    friend_requester_id IN
    (SELECT id FROM users
      WHERE username = "reese");
```

```
DELETE FROM users WHERE username = "reese";
```

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

SQLite documentation from www.sqlite.org, and
Example queries from www.sqlitetutorial.net.
Sqlite3 database.

2. About how long, in minutes, did you spend on this problem's questions?

180 minutes.

Less is More

1. The given sequence can be encoded using run-length encoding as: **T3A4C2G1A3**.
2. When the original DNA sequence has each different nucleotide appearing only once. In other words, there are no consecutive occurrences of the same nucleotide in a DNA sequence. For example, the following DNA sequence takes up 6 characters:

ATGCTA

However, its run-length encoding takes up 12 characters:

A1T1G1C1T1A1

3. Flag of Germany could be compressed more. The reason is that the width of the three colors in the flag of Germany is much larger than the width of the three colors in the flag of Romania. Therefore, the number of pixel rows for each color in the flag of Germany is much smaller than that of Romania. Since our compression works by storing just one pixel data followed by the count of the pixels, the flag of Germany will need a smaller number of storage than the flag of Romania.

4. Using the proposed encoding method, **CCCAGTTA** can be encoded as,

11101011110

5. Using the proposed encoding method, **TGCATCCA** can be encoded as,

11101011110

6. We note that, using the proposed encoding method, we are obtaining identical encoding for different DNA sequences. For example, the nucleotide sequences **CCCAGTTA** and **TGCATCCA** above in Part 4 and 5 result in the same encoding: **11101011110**

We can improve the above encoding technique by using two bits for each nucleotide. In this new technique, we will use the following encoding:

- A is represented as 00
- C is represented as 01
- G is represented as 10
- T is represented as 11

During decoding, the compressed sequence will be decoded from the left end two bits at a time using the above mapping. Therefore, the above encoding will result in unique encoding for distinct DNA sequences. Since each nucleotide takes up only 2 bits as opposed to 8 bits using ASCII, our new approach will use significantly fewer bits.

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

None

2. About how long, in minutes, did you spend on this problem's questions?

30 minutes

Random Questions

1. The average running time of **sorted()** is $O(n \log n)$. Since the comparison of two lists on the average takes $O(n)$ time and the **issorted()** function compares two lists using the `==` operator, the overall average running time of **issorted()** is $O(n^2 \log n)$.
2. The following implementation of **issorted()** has a running time of $O(n)$, as it is comparing each element in the list with the previous element. On the average, it will examine $n/2$ elements, and in the worst-case, it will examine the entire list.

```
def issorted(numbers):
    size_of_list = len(numbers)

    for i in range(1, size_of_list):
        if numbers[i-1] > numbers[i]:
            return False
    return True
```

3. The running time of **issorted()** cannot ever be $O(1)$, which means a constant number of operations. The reason is that on the average, **issorted()** will need to do $n/2$ comparisons of adjacent elements, and in the worst-case, it will have to compare the adjacent elements in the entire list. In the best-case scenario, numbers is a sorted list. In order to verify that it is a sorted list, we must iterate through the list once, which takes $O(n)$ time.
4. In the worst case, the running time of sort (called Bogo sort) will be $O(\text{infinity})$ as the algorithm does not have an upper bound. However, the average running time will be in the order of $O(n \times n!)$. The reason is that there is a total of $n!$ different ordering for the input list, and on the average shuffling finds the sorted list in $O(n!)$ time. The shuffling itself takes $O(n)$ time.
5. In the best-case scenario, **itertools.permutations()** returns the list of numbers in sorted order immediately in $O(1)$ time. Therefore, considering the time spent in **issorted()** function, the overall lower bound on the running time is $O(n)$.
6. In the worst-case scenario, **itertools.permutations()** returns the list of numbers in sorted order only on its last iteration. This is $O(n!)$. Therefore, considering the time spent in **issorted()** function, the upper bound on the running time is $O(n \times n!)$.

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?
Bogosort on Wikipedia
2. About how long, in minutes, did you spend on this problem's questions?
30 minutes

Teetering on the Edge

1. Yes, because we can place a fulcrum between both elements and their values are equal.
2. No, because the sum of the elements is odd. So, we cannot separate the elements equally on both sides.
3. Yes. We can place the elements [3, 4, 6] on the left side of the fulcrum, and the remaining elements [5, 8] on the right side of the fulcrum.
4. The first property is the sum of the elements. If the sum of the elements is odd, then the elements cannot be balanced as we cannot divide their values equally on both sides of the fulcrum. Assuming that the sum of the elements is even, the second property is that if the value of one element is more than half of the values of all the remaining elements, the list is not balanceable.
5. Please find my implementation of balanceable on the next page. *I have submitted my code for this question via `submit50`.*

```

def balanceable(numbers):
    import itertools

    # Find the sum of all the elements in the list
    total = sum(numbers)

    # If the sum is odd, the list cannot be balanced
    if (total % 2 == 1):
        return False

    # Find half of the above sum
    half_of_sum = total//2

    # If the largest value in the list is larger than
    # half of the total, then the list cannot be balanced
    if (max(numbers) > half_of_sum):
        return False

    # Let us consider each of the n! permutations for the
    # number list.
    for permutation in itertools.permutations(numbers):
        sum_so_far = 0
        no_of_elements = len(numbers)

        for i in range(no_of_elements):
            sum_so_far = sum_so_far + permutation[i]

            # If sum_so_far exceeds half_of_sum, no need to check
            # further.
            if (sum_so_far >= half_of_sum):
                break

        # If sum_so_far is exactly equal to half of the sum of the list,
        # then the list is balanceable
        if (sum_so_far == half_of_sum):
            return True

    # If we reach here, we have examined all n! permutations of
    # the number list and confirmed that the list is not balanceable
    #
    return False

```

Debrief

1. Which resources, if any, did you find helpful in answering this problem's questions?

Python tutorials on www.tutorialpoint.com

2. About how long, in minutes, did you spend on this problem's questions?

60 minutes